

F2FS in-mem data structures

fs/f2fs/f2fs.h

```
733 enum {
734     CP_TIME,
735     REQ_TIME,
736     MAX_TIME,
737 };
738
739 #ifdef CONFIG_F2FS_FS_ENCRYPTION
740 #define F2FS_KEY_DESC_PREFIX "f2fs:"
741 #define F2FS_KEY_DESC_PREFIX_SIZE 5
742 #endif
743 struct f2fs_sb_info {
744     struct super_block *sb;          /* pointer to VFS super block */
745     struct proc_dir_entry *s_proc;    /* proc entry */
746     struct f2fs_super_block *raw_super; /* raw super block pointer */
747     int valid_super_block;            /* valid super block no */
748     unsigned long s_flag;             /* flags for sbi */
749
750 #ifdef CONFIG_F2FS_FS_ENCRYPTION
751     u8 key_prefix[F2FS_KEY_DESC_PREFIX_SIZE];
752     u8 key_prefix_size;
753 #endif
754     /* for node-related operations */
755     struct f2fs_nm_info *nm_info;     /* node manager */
756     struct inode *node_inode;         /* cache node blocks */
757
758     /* for segment-related operations */
759     struct f2fs_sm_info *sm_info;     /* segment manager */
760
761     /* for bio operations */
762     struct f2fs_bio_info read_io;     /* for read bios */
763     struct f2fs_bio_info write_io[NR_PAGE_TYPE]; /* for write bios */
764     struct mutex wio_mutex[NODE + 1]; /* bio ordering for NODE/DATA */
765
766     /* for checkpoint */
767     struct f2fs_checkpoint *ckpt;     /* raw checkpoint pointer */
768     int cur_cp_pack;                  /* remain current cp pack */
769     spinlock_t cp_lock;               /* for flag in ckpt */
770     struct inode *meta_inode;         /* cache meta blocks */
771     struct mutex cp_mutex;            /* checkpoint procedure lock */
772     struct rw_semaphore cp_rwsem;     /* blocking FS operations */
773     struct rw_semaphore node_write;   /* locking node writes */
774     wait_queue_head_t cp_wait;
775     unsigned long last_time[MAX_TIME]; /* to store time in jiffies */
776     long interval_time[MAX_TIME];     /* to store thresholds */
777
778     struct inode_management im[MAX_INO_ENTRY]; /* manage inode cache */
779
780     /* for orphan inode, use 0th array */
781     unsigned int max_orphans;         /* max orphan inodes */
782
783     /* for inode management */
784     struct list_head inode_list[NR_INODE_TYPE]; /* dirty inode list */
785     spinlock_t inode_lock[NR_INODE_TYPE]; /* for dirty inode list lock */
786
787     /* for extent tree cache */
788     struct radix_tree_root extent_tree_root; /* cache extent cache entries */
789     struct rw_semaphore extent_tree_lock; /* locking extent radix tree */
790     struct list_head extent_list; /* lru list for shrinker */
791     spinlock_t extent_lock; /* locking extent lru list */
792     atomic_t total_ext_tree; /* extent tree count */
793     struct list_head zombie_list; /* extent zombie tree list */
794     atomic_t total_zombie_tree; /* extent zombie tree count */

```

```

795 atomic_t total_ext_node;          /* extent info count */
796
797 /* basic filesystem units */
798 unsigned int log_sectors_per_block; /* log2 sectors per block */
799 unsigned int log_blocksize;        /* log2 block size */
800 unsigned int blocksize;            /* block size */
801 unsigned int root_ino_num;         /* root inode number */
802 unsigned int node_ino_num;         /* node inode number */
803 unsigned int meta_ino_num;         /* meta inode number */
804 unsigned int log_blocks_per_seg;   /* log2 blocks per segment */
805 unsigned int blocks_per_seg;      /* blocks per segment */
806 unsigned int segs_per_sec;         /* segments per section */
807 unsigned int secs_per_zone;        /* sections per zone */
808 unsigned int total_sections;       /* total section count */
809 unsigned int total_node_count;     /* total node block count */
810 unsigned int total_valid_node_count; /* valid node block count */
811 loff_t max_file_blocks;            /* max block index of file */
812 int active_logs;                   /* # of active logs */
813 int dir_level;                     /* directory level */
814
815 block_t user_block_count;          /* # of user blocks */
816 block_t total_valid_block_count;   /* # of valid blocks */
817 block_t discard_blks;              /* discard command candidates */
818 block_t last_valid_block_count;    /* for recovery */
819 u32 s_next_generation;             /* for NFS support */
820 atomic_t nr_wb_bios;               /* # of writeback bios */
821
822 /* # of pages, see count_type */
823 atomic_t nr_pages[NR_COUNT_TYPE];
824 /* # of allocated blocks */
825 struct percpu_counter alloc_valid_block_count;
826
827 /* valid inode count */
828 struct percpu_counter total_valid_inode_count;
829
830 struct f2fs_mount_info mount_opt; /* mount options */
831
832 /* for cleaning operations */
833 struct mutex gc_mutex;             /* mutex for GC */
834 struct f2fs_gc_kthread *gc_thread; /* GC thread */
835 unsigned int cur_victim_sec;        /* current victim section num */
836
837 /* threshold for converting bg victims for fg */
838 u64 fggc_threshold;
839
840 /* maximum # of trials to find a victim segment for SSR and GC */
841 unsigned int max_victim_search;
842
843 /*
844  * for stat information.
845  * one is for the LFS mode, and the other is for the SSR mode.
846  */
847 #ifdef CONFIG_F2FS_STAT_FS
848 struct f2fs_stat_info *stat_info; /* FS status information */
849 unsigned int segment_count[2];    /* # of allocated segments */
850 unsigned int block_count[2];      /* # of allocated blocks */
851 atomic_t inplace_count;           /* # of inplace update */
852 atomic64_t total_hit_ext;         /* # of lookup extent cache */
853 atomic64_t read_hit_rbtrees;     /* # of hit rbtrees extent node */
854 atomic64_t read_hit_largest;     /* # of hit largest extent node */
855 atomic64_t read_hit_cached;      /* # of hit cached extent node */
856 atomic_t inline_xattr;            /* # of inline_xattr inodes */
857 atomic_t inline_inode;            /* # of inline_data inodes */
858 atomic_t inline_dir;              /* # of inline_dentry inodes */
859 int bg_gc;                         /* background gc calls */
860 unsigned int ndirty_inode[NR_INODE_TYPE]; /* # of dirty inodes */
861 #endif
862 unsigned int last_victim[2];       /* last victim segment # */
863 spinlock_t stat_lock;              /* lock for stat operations */

```

```

864 spinlock_t stat_lock; /* lock for stat operations */
865 /* For sysfs support */
866 struct kobject s_kobj;
867 struct completion s_kobj_unregister;
868
869 /* For shrinker support */
870 struct list_head s_list;
871 struct mutex umount_mutex;
872 unsigned int shrinker_run_no;
873
874 /* For write statistics */
875 u64 sectors_written_start;
876 u64 kbytes_written;
877
878 /* Reference to checksum algorithm driver via cryptoapi */
879 struct crypto_shash *s_chksum_driver;
880
881 /* For fault injection */
882 #ifdef CONFIG_F2FS_FAULT_INJECTION
883 struct f2fs_fault_info fault_info;
884 #endif
885 };

```

```

606 struct f2fs_sm_info {
607     struct sit_info *sit_info; /* whole segment information */
608     struct free_segmap_info *free_info; /* free segment information */
609     struct dirty_seglist_info *dirty_info; /* dirty segment information */
610     struct curseg_info *curseg_array; /* active segment information */
611
612     block_t seg0_blkaddr; /* block address of 0'th segment */
613     block_t main_blkaddr; /* start block address of main area */
614     block_t ssa_blkaddr; /* start block address of SSA area */
615
616     unsigned int segment_count; /* total # of segments */
617     unsigned int main_segments; /* # of segments in main area */
618     unsigned int reserved_segments; /* # of reserved segments */
619     unsigned int ovp_segments; /* # of overprovision segments */
620
621     /* a threshold to reclaim prefree segments */
622     unsigned int rec_prefree_segments;
623
624     /* for small discard management */
625     struct list_head discard_list; /* 4KB discard list */
626     struct list_head wait_list; /* linked with issued discard bio */
627     int nr_discards; /* # of discards in the list */
628     int max_discards; /* max. discards to be issued */
629
630     /* for batched trimming */
631     unsigned int trim_sections; /* # of sections to trim */
632
633     struct list_head sit_entry_set; /* sit entry set list */
634
635     unsigned int ipu_policy; /* in-place-update policy */
636     unsigned int min_ipu_util; /* in-place-update threshold */
637     unsigned int min_fsync_blocks; /* threshold for fsync */
638
639     /* for flush command control */
640     struct flush_cmd_control *cmd_control_info;
641
642 };

```

```

507 struct f2fs_nm_info {
508     block_t nat_blkaddr;      /* base disk address of NAT */
509     nid_t max_nid;            /* maximum possible node ids */
510     nid_t available_nids;     /* maximum available node ids */
511     nid_t next_scan_nid;      /* the next nid to be scanned */
512     unsigned int ram_thresh;   /* control the memory footprint */
513     unsigned int ra_nid_pages; /* # of nid pages to be readahead */
514     unsigned int dirty_nats_ratio; /* control dirty nats ratio threshold */
515
516     /* NAT cache management */
517     struct radix_tree_root nat_root; /* root of the nat entry cache */
518     struct radix_tree_root nat_set_root; /* root of the nat set cache */
519     struct rw_semaphore nat_tree_lock; /* protect nat_tree_lock */
520     struct list_head nat_entries; /* cached nat entry list (clean) */
521     unsigned int nat_cnt; /* the # of cached nat entries */
522     unsigned int dirty_nat_cnt; /* total num of nat entries in set */
523
524     /* free node ids management */
525     struct radix_tree_root free_nid_root; /* root of the free_nid cache */
526     struct list_head free_nid_list; /* a list for free nids */
527     spinlock_t free_nid_list_lock; /* protect free nid list */
528     unsigned int fcmt; /* the number of free node id */
529     struct mutex build_lock; /* lock for build free nids */
530
531     /* for checkpoint */
532     char *nat_bitmap; /* NAT bitmap pointer */
533     int bitmap_size; /* bitmap size */
534 };

```

```

419 struct f2fs_inode_info {
420     struct inode vfs_inode; /* serve a vfs inode */
421     unsigned long i_flags; /* keep an inode flags for ioctl */
422     unsigned char i_advise; /* use to give file attribute hints */
423     unsigned char i_dir_level; /* use for dentry level for large dir */
424     unsigned int i_current_depth; /* use only in directory structure */
425     unsigned int i_pino; /* parent inode number */
426     umode_t i_acl_mode; /* keep file acl mode temporarily */
427
428     /* Use below internally in f2fs */
429     unsigned long flags; /* use to pass per-file flags */
430     struct rw_semaphore i_sem; /* protect fi info */
431     atomic_t dirty_pages; /* # of dirty pages */
432     f2fs_hash_t chash; /* hash value of given file name */
433     unsigned int clevel; /* maximum level of given file name */
434     struct task_struct *task; /* lookup and create consistency */
435     nid_t i_xattr_nid; /* node id that contains xattrs */
436     unsigned long long xattr_ver; /* cp version of xattr modification */
437     loff_t last_disk_size; /* lastly written file size */
438
439     struct list_head dirty_list; /* dirty list for dirs and files */
440     struct list_head gdirty_list; /* linked in global dirty list */
441     struct list_head inmem_pages; /* inmemory pages managed by f2fs */
442     struct mutex inmem_lock; /* lock for inmemory pages */
443     struct extent_tree *extent_tree; /* cached extent_tree entry */
444     struct rw_semaphore dio_rwsem[2]; /* avoid racing between dio and gc */
445 };

```

```

286 /*
287  * For INODE and NODE manager
288  */
289 /* for directory operations */
290 struct f2fs_dentry_ptr {

```

```

291 struct inode *inode;
292 const void *bitmap;
293 struct f2fs_dir_entry *dentry;
294 __u8 (*filename)[F2FS_SLOT_LEN];
295 int max;
296 };

```

```

663 /*
664  * The below are the page types of bios used in submit_bio().
665  * The available types are:
666  * DATA          User data pages. It operates as async mode.
667  * NODE           Node pages. It operates as async mode.
668  * META           FS metadata pages such as SIT, NAT, CP.
669  * NR_PAGE_TYPE   The number of page types.
670  * META_FLUSH     Make sure the previous pages are written
671  *                with waiting the bio's completion
672  * ...            Only can be used with META.
673  */
674 #define PAGE_TYPE_OF_BIO(type) ((type) > META ? META : (type))
675 enum page_type {
676     DATA,
677     NODE,
678     META,
679     NR_PAGE_TYPE,
680     META_FLUSH,
681     INMEM, /* the below types are used by tracepoints only. */
682     INMEM_DROP,
683     INMEM_REVOKE,
684     IPU,
685     OPU,
686 };

```

```

708 enum inode_type {
709     DIR_INODE, /* for dirty dir inode */
710     FILE_INODE, /* for dirty regular/symlink inode */
711     DIRTY_META, /* for all dirtied inode metadata */
712     NR_INODE_TYPE,
713 };

715 /* for inner inode cache management */
716 struct inode_management {
717     struct radix_tree_root ino_root; /* ino entry array */
718     spinlock_t ino_lock; /* for ino entry lock */
719     struct list_head ino_list; /* inode list head */
720     unsigned long ino_num; /* number of entries */
721 };
722
723 /* For s_flag in struct f2fs_sb_info */

```

```

644 /*
645  * For superblock
646 */
647 /*
648  * COUNT_TYPE for monitoring
649  *
650  * f2fs monitors the number of several block types such as on-writeback,
651  * dirty dentry blocks, dirty node blocks, and dirty meta blocks.
652 */
653 enum count_type {
654     F2FS_DIRTY_DENTS,
655     F2FS_DIRTY_DATA,
656     F2FS_DIRTY_NODES,
657     F2FS_DIRTY_META,
658     F2FS_INMEM_PAGES,
659     F2FS_DIRTY_IMETA,
660     NR_COUNT_TYPE,
661 };

```

```

688 struct f2fs_io_info {
689     struct f2fs_sb_info *sbi; /* f2fs_sb_info pointer */
690     enum page_type type; /* contains DATA/NODE/META/META_FLUSH */
691     int op; /* contains REQ_OP_ */
692     int op_flags; /* rq_flag_bits */
693     block_t new_blkaddr; /* new block address to be written */
694     block_t old_blkaddr; /* old block address before Cow */
695     struct page *page; /* page to be written */
696     struct page *encrypted_page; /* encrypted page */
697 };
698
699 #define is_read_io(rw) (rw == READ)
700 struct f2fs_bio_info {
701     struct f2fs_sb_info *sbi; /* f2fs superblock */
702     struct bio *bio; /* bios to merge */
703     sector_t last_block_in_bio; /* last block number */
704     struct f2fs_io_info fio; /* store buffered io info. */
705     struct rw_semaphore io_rwsem; /* blocking op for bio */
706 };

```

```

564 /*
565  * For SIT manager
566  *
567  * By default, there are 6 active log areas across the whole main area.
568  * When considering hot and cold data separation to reduce cleaning overhead,
569  * we split 3 for data logs and 3 for node logs as hot, warm, and cold types,
570  * respectively.
571  * In the current design, you should not change the numbers intentionally.
572  * Instead, as a mount option such as active_logs=x, you can use 2, 4, and 6
573  * logs individually according to the underlying devices. (default: 6)
574  * Just in case, on-disk layout covers maximum 16 logs that consist of 8 for
575  * data and 8 for node logs.
576 */
577 #define NR_CURSEG_DATA_TYPE (3)
578 #define NR_CURSEG_NODE_TYPE (3)
579 #define NR_CURSEG_TYPE (NR_CURSEG_DATA_TYPE +
NR_CURSEG_NODE_TYPE)
580

```

```

581 enum {
582     CURSEG_HOT_DATA = 0, /* directory entry blocks */
583     CURSEG_WARM_DATA, /* data blocks */
584     CURSEG_COLD_DATA, /* multimedia or GCed data blocks */
585     CURSEG_HOT_NODE, /* direct node blocks of directory files */
586     CURSEG_WARM_NODE, /* direct node blocks of normal files */
587     CURSEG_COLD_NODE, /* indirect node blocks */
588     NO_CHECK_TYPE,
589     CURSEG_DIRECT_IO, /* to use for the direct IO path */
590 };

```

```

592 struct flush_cmd {
593     struct completion wait;
594     struct llist_node llnode;
595     int ret;
596 };
597
598 struct flush_cmd_control {
599     struct task_struct *f2fs_issue_flush; /* flush thread */
600     wait_queue_head_t flush_wait_queue; /* waiting queue for wake-up */
601     atomic_t submit_flush; /* # of issued flushes */
602     struct llist_head issue_list; /* list for command issue */
603     struct llist_node *dispatch_list; /* list for command dispatch */
604 };

```

```

536 /*
537  * this structure is used as one of function parameters.
538  * all the information are dedicated to a given direct node block determined
539  * by the data offset in a file.
540  */
541 struct dnode_of_data {
542     struct inode *inode; /* vfs inode pointer */
543     struct page *inode_page; /* its inode page, NULL is possible */
544     struct page *node_page; /* cached direct node page */
545     nid_t nid; /* node id of the direct node block */
546     unsigned int ofs_in_node; /* data offset in the node page */
547     bool inode_page_locked; /* inode page is locked or not */
548     bool node_changed; /* is node block changed */
549     char cur_level; /* level of hole node page */
550     char max_level; /* level of current page located */
551     block_t data_blkaddr; /* block address of the node block */
552 };

```

```

370 /*
371  * This structure is taken from ext4_map_blocks.
372  *
373  * Note that, however, f2fs uses NEW and MAPPED flags for f2fs_map_blocks().
374  */
375 #define F2FS_MAP_NEW (1 << BH_New)
376 #define F2FS_MAP_MAPPED (1 << BH_Mapped)
377 #define F2FS_MAP_UNWRITTEN (1 << BH_Unwritten)
378 #define F2FS_MAP_FLAGS (F2FS_MAP_NEW | F2FS_MAP_MAPPED |
379     F2FS_MAP_UNWRITTEN)
380
381 struct f2fs_map_blocks {
382     block_t m_pblk;
383     block_t m_lblk;

```



```

384     unsigned int m_len;
385     unsigned int m_flags;
386     pgoff_t *m_next_pgofs;    /* point next possible non-hole pgofs */
387 };
388
389 /* for flag in get_data_block */
390 #define F2FS_GET_BLOCK_READ      0
391 #define F2FS_GET_BLOCK_DIO      1
392 #define F2FS_GET_BLOCK_FIEMAP   2
393 #define F2FS_GET_BLOCK_BMAP     3
394 #define F2FS_GET_BLOCK_PRE_DIO  4
395 #define F2FS_GET_BLOCK_PRE_AIO  5
396
397 /*
398  * i_advise uses FADVISE_XXX_BIT. We can add additional hints later.
399  */
400 #define FADVISE_COLD_BIT        0x01
401 #define FADVISE_LOST_PINO_BIT   0x02
402 #define FADVISE_ENCRYPT_BIT     0x04
403 #define FADVISE_ENC_NAME_BIT   0x08
404
405 #define file_is_cold(inode)     is_file(inode, FADVISE_COLD_BIT)
406 #define file_wrong_pino(inode) is_file(inode, FADVISE_LOST_PINO_BIT)
407 #define file_set_cold(inode)    set_file(inode, FADVISE_COLD_BIT)
408 #define file_lost_pino(inode)   set_file(inode, FADVISE_LOST_PINO_BIT)
409 #define file_clear_cold(inode)  clear_file(inode, FADVISE_COLD_BIT)
410 #define file_got_pino(inode)    clear_file(inode, FADVISE_LOST_PINO_BIT)
411 #define file_is_encrypt(inode)  is_file(inode, FADVISE_ENCRYPT_BIT)
412 #define file_set_encrypt(inode) set_file(inode, FADVISE_ENCRYPT_BIT)
413 #define file_clear_encrypt(inode) clear_file(inode, FADVISE_ENCRYPT_BIT)
414 #define file_enc_name(inode)    is_file(inode, FADVISE_ENC_NAME_BIT)
415 #define file_set_enc_name(inode) set_file(inode, FADVISE_ENC_NAME_BIT)
416
417 #define DEF_DIR_LEVEL          0

```

```

318 /*
319  * XATTR_NODE_OFFSET stores xattrs to one node block per file keeping -1
320  * as its node offset to distinguish from index node blocks.
321  * But some bits are used to mark the node block.
322  */
323 #define XATTR_NODE_OFFSET      (((unsigned int)-1) << OFFSET_BIT_SHIFT) \
324                                >> OFFSET_BIT_SHIFT)
325 enum {
326     ALLOC_NODE,          /* allocate a new node page if needed */
327     LOOKUP_NODE,         /* look up a node without readahead */
328     LOOKUP_NODE_RA,      /*
329                          * look up a node with readahead called
330                          * by get_data_block.
331                          */
332 };
333
334 #define F2FS_LINK_MAX  0xffffffff /* maximum link count per file */
335
336 #define MAX_DIR_RA_PAGES  4 /* maximum ra pages of dir */
337
338 /* vector size for gang look-up from extent cache that consists of radix tree */
339 #define EXT_TREE_VEC_SIZE  64
340
341 /* for in-memory extent cache entry */
342 #define F2FS_MIN_EXTENT_LEN  64 /* minimum extent length */
343
344 /* number of extent info in extent cache we try to shrink */
345 #define EXTENT_CACHE_SHRINK_NUMBER  128
346

```



```

347 struct extent_info {
348     unsigned int fofs;          /* start offset in a file */
349     u32 blk;                    /* start block address of the extent */
350     unsigned int len;           /* length of the extent */
351 };
352
353 struct extent_node {
354     struct rb_node rb_node;     /* rb node located in rb-tree */
355     struct list_head list;      /* node in global extent list of sbi */
356     struct extent_info ei;      /* extent info */
357     struct extent_tree *et;     /* extent tree pointer */
358 };
359
360 struct extent_tree {
361     nid_t ino;                  /* inode number */
362     struct rb_root root;       /* root of extent info rb-tree */
363     struct extent_node *cached_en; /* recently accessed extent node */
364     struct extent_info largest; /* largest extent info */
365     struct list_head list;      /* to be used by sbi->zombie_list */
366     rwlock_t lock;             /* protect extent info rb-tree */
367     atomic_t node_cnt;         /* # of extent node in rb-tree */
368 };

```

```

235 #define F2FS_IOC_GETFLAGS      FS_IOC_GETFLAGS
236 #define F2FS_IOC_SETFLAGS      FS_IOC_SETFLAGS
237 #define F2FS_IOC_GETVERSION    FS_IOC_GETVERSION
238
239 #define F2FS_IOCTL_MAGIC       0xf5
240 #define F2FS_IOC_START_ATOMIC_WRITE  _IO(F2FS_IOCTL_MAGIC, 1)
241 #define F2FS_IOC_COMMIT_ATOMIC_WRITE  _IO(F2FS_IOCTL_MAGIC, 2)
242 #define F2FS_IOC_START_VOLATILE_WRITE  _IO(F2FS_IOCTL_MAGIC, 3)
243 #define F2FS_IOC_RELEASE_VOLATILE_WRITE _IO(F2FS_IOCTL_MAGIC, 4)
244 #define F2FS_IOC_ABORT_VOLATILE_WRITE  _IO(F2FS_IOCTL_MAGIC, 5)
245 #define F2FS_IOC_GARBAGE_COLLECT       _IO(F2FS_IOCTL_MAGIC, 6)
246 #define F2FS_IOC_WRITE_CHECKPOINT      _IO(F2FS_IOCTL_MAGIC, 7)
247 #define F2FS_IOC_DEFRAGMENT            _IO(F2FS_IOCTL_MAGIC, 8)
248 #define F2FS_IOC_MOVE_RANGE             _IOWR(F2FS_IOCTL_MAGIC, 9, \
249         struct f2fs_move_range)
250
251 #define F2FS_IOC_SET_ENCRYPTION_POLICY  FS_IOC_SET_ENCRYPTION_POLICY
252 #define F2FS_IOC_GET_ENCRYPTION_POLICY  FS_IOC_GET_ENCRYPTION_POLICY
253 #define F2FS_IOC_GET_ENCRYPTION_PWSALT  FS_IOC_GET_ENCRYPTION_PWSALT
254
255 /*
256  * should be same as XFS_IOC_GOINGDOWN.
257  * Flags for going down operation used by FS_IOC_GOINGDOWN
258  */
259 #define F2FS_IOC_SHUTDOWN      _IOR('X', 125, __u32) /* Shutdown */
260 #define F2FS_GOING_DOWN_FULLSYNC  0x0 /* going down with full sync */
261 #define F2FS_GOING_DOWN_METASYNC  0x1 /* going down with metadata
*/
262 #define F2FS_GOING_DOWN_NOSYNC    0x2 /* going down */
263 #define F2FS_GOING_DOWN_METAFLUSH 0x3 /* going down with meta flush
*/
264
265 #if defined(__KERNEL__) && defined(CONFIG_COMPAT)
266 /*
267  * ioctl commands in 32 bit emulation
268  */
269 #define F2FS_IOC32_GETFLAGS      FS_IOC32_GETFLAGS
270 #define F2FS_IOC32_SETFLAGS      FS_IOC32_SETFLAGS
271 #define F2FS_IOC32_GETVERSION    FS_IOC32_GETVERSION
272 #endif

```

```

274 struct f2fs_defragment {
275     u64 start;
276     u64 len;
277 };
278
279 struct f2fs_move_range {
280     u32 dst_fd;          /* destination fd */
281     u64 pos_in;          /* start position in src_fd */
282     u64 pos_out;         /* start position in dst_fd */
283     u64 len;             /* size to move */
284 };

```

```

139 struct cp_control {
140     int reason;
141     __u64 trim_start;
142     __u64 trim_end;
143     __u64 trim_minlen;
144     __u64 trimmed;
145 };
146
147 /*
148  * For CP/NAT/SIT/SSA readahead
149  */
150 enum {
151     META_CP,
152     META_NAT,
153     META_SIT,
154     META_SSA,
155     META_POR,
156 };
157
158 /* for the list of ino */
159 enum {
160     ORPHAN_INO,          /* for orphan ino list */
161     APPEND_INO,          /* for append ino list */
162     UPDATE_INO,          /* for update ino list */
163     MAX_INO_ENTRY,       /* max. list */
164 };
165
166 struct ino_entry {
167     struct list_head list; /* list head */
168     nid_t ino;             /* inode number */
169 };
170
171 /* for the list of inodes to be GCed */
172 struct inode_entry {
173     struct list_head list; /* list head */
174     struct inode *inode;   /* vfs inode pointer */
175 };
176
177 /* for the list of blockaddresses to be discarded */
178 struct discard_entry {
179     struct list_head list; /* list head */
180     block_t blkaddr;       /* block address to be discarded */
181     int len;               /* # of consecutive blocks of the discard */
182 };
183
184 struct bio_entry {
185     struct list_head list;
186     struct bio *bio;
187     struct completion event;
188     int error;
189 };
190
191 /* for the list of fsync inodes, used only during recovery */
192 struct fsync_inode_entry {
193     struct list_head list; /* list head */

```

```
194     struct inode *inode; /* vfs inode pointer */
195     block_t blkaddr; /* block address locating the last fsync */
196     block_t last_dentry; /* block address locating the last dentry */
197 };
```

```
25 struct f2fs_gc_kthread {
26     struct task_struct *f2fs_gc_task;
27     wait_queue_head_t gc_wait_queue_head;
28
29     /* for gc sleep time */
30     unsigned int min_sleep_time;
31     unsigned int max_sleep_time;
32     unsigned int no_gc_sleep_time;
33
34     /* for changing gc mode */
35     unsigned int gc_idle;
36 };
37
38 struct gc_inode_list {
39     struct list_head ilist;
40     struct radix_tree_root iroot;
41 };
```

fs/f2fs/node.h

```
41 /* For flag in struct node_info */
42 enum {
43     IS_CHECKPOINTED,    /* is it checkpointed before? */
44     HAS_FSYNCHED_INODE, /* is the inode fsynced before? */
45     HAS_LAST_FSYNC,     /* has the latest node fsync mark? */
46     IS_DIRTY,           /* this nat entry is dirty? */
47 };
48
49 /*
50  * For node information
51  */
52 struct node_info {
53     nid_t nid;           /* node id */
54     nid_t ino;           /* inode number of the node's owner */
55     block_t blk_addr;    /* block address of the node */
56     unsigned char version; /* version of the node */
57     unsigned char flag;   /* for node information bits */
58 };
59
60 struct nat_entry {
61     struct list_head list; /* for clean or dirty nat list */
62     struct node_info ni;   /* in-memory node information */
63 };

137 enum mem_type {
138     FREE_NIDS,    /* indicates the free nid list */
139     NAT_ENTRIES,  /* indicates the cached nat entry */
140     DIRTY_DENTS,  /* indicates dirty dentry pages */
141     INO_ENTRIES,  /* indicates inode entries */
142     EXTENT_CACHE, /* indicates extent cache */
143     BASE_CHECK,   /* check kernel status */
144 };
145
146 struct nat_entry_set {
147     struct list_head set_list; /* link with other nat sets */
148     struct list_head entry_list; /* link with dirty nat entries */
149     nid_t set;                 /* set number */
150     unsigned int entry_cnt;     /* the # of nat entries in set */
151 };
152
153 /*
154  * For free nid mangement
155  */
156 enum nid_state {
157     NID_NEW,    /* newly added to free nid list */
158     NID_ALLOC   /* it is allocated */
159 };
160
161 struct free_nid {
162     struct list_head list; /* for free node id list */
163     nid_t nid;             /* node id */
164     int state;             /* in use or not: NID_NEW or NID_ALLOC */
165 };
```

fs/f2fs/segment.h

```
149 /* for a function parameter to select a victim segment */
150 struct victim_sel_policy {
151     int alloc_mode;          /* LFS or SSR */
152     int gc_mode;             /* GC_CB or GC_GREEDY */
153     unsigned long *dirty_segmap; /* dirty segment bitmap */
154     unsigned int max_search; /* maximum # of segments to search */
155     unsigned int offset;     /* last scanned bitmap offset */
156     unsigned int ofs_unit;   /* bitmap search unit */
157     unsigned int min_cost;   /* minimum cost */
158     unsigned int min_segno;  /* segment # having min. cost */
159 };
160
161 struct seg_entry {
162     unsigned int type:6;      /* segment type like CURSEG_XXX_TYPE */
163     unsigned int valid_blocks:10; /* # of valid blocks */
164     unsigned int ckpt_valid_blocks:10; /* # of valid blocks last cp */
165     unsigned int padding:6;    /* padding */
166     unsigned char *cur_valid_map; /* validity bitmap of blocks */
167     /*
168      * # of valid blocks and the validity bitmap stored in the the last
169      * checkpoint pack. This information is used by the SSR mode.
170      */
171     unsigned char *ckpt_valid_map; /* validity bitmap of blocks last cp */
172     unsigned char *discard_map;
173     unsigned long long mtime; /* modification time of the segment */
174 };
175
176 struct sec_entry {
177     unsigned int valid_blocks; /* # of valid blocks in a section */
178 };
179
180 struct segment_allocation {
181     void (*allocate_segment)(struct f2fs_sb_info *, int, bool);
182 };
183
184 struct inmem_pages {
185     struct list_head list;
186     struct page *page;
187     block_t old_addr; /* for revoking when fail to commit */
188 };
189
190 struct sit_info {
191     const struct segment_allocation *s_ops;
192
193     block_t sit_base_addr; /* start block address of SIT area */
194     block_t sit_blocks; /* # of blocks used by SIT area */
195     block_t written_valid_blocks; /* # of valid blocks in main area */
196     char *sit_bitmap; /* SIT bitmap pointer */
197     unsigned int bitmap_size; /* SIT bitmap size */
198
199     unsigned long *tmp_map; /* bitmap for temporal use */
200     unsigned long *dirty_sentries_bitmap; /* bitmap for dirty sentries */
201     unsigned int dirty_sentries; /* # of dirty sentries */
202     unsigned int sents_per_block; /* # of SIT entries per block */
203     struct mutex sentry_lock; /* to protect SIT cache */
204     struct seg_entry *sentries; /* SIT segment-level cache */
205     struct sec_entry *sec_entries; /* SIT section-level cache */
206
207     /* for cost-benefit algorithm in cleaning procedure */
208     unsigned long long elapsed_time; /* elapsed time after mount */
209     unsigned long long mounted_time; /* mount time */
210     unsigned long long min_mtime; /* min. modification time */
211     unsigned long long max_mtime; /* max. modification time */
212 };
213
```

fs/f2fs/segment.h

223 struct free_segmap_info {

```
224     unsigned int start_segno;    /* start segment number logically */
225     unsigned int free_segments;   /* # of free segments */
226     unsigned int free_sections;   /* # of free sections */
227     spinlock_t segmap_lock;      /* free segmap lock */
228     unsigned long *free_segmap;   /* free segment bitmap */
229     unsigned long *free_secmap;   /* free section bitmap */
```

```
230 };
```

```
231
```

```
232 /* Notice: The order of dirty type is same with CURSEG_XXX in f2fs.h */
```

233 enum dirty_type {

```
234     DIRTY_HOT_DATA,             /* dirty segments assigned as hot data logs */
235     DIRTY_WARM_DATA,            /* dirty segments assigned as warm data logs */
236     DIRTY_COLD_DATA,           /* dirty segments assigned as cold data logs */
237     DIRTY_HOT_NODE,            /* dirty segments assigned as hot node logs */
238     DIRTY_WARM_NODE,           /* dirty segments assigned as warm node logs */
```

```
239     DIRTY_COLD_NODE,          /* dirty segments assigned as cold node logs */
```

```
240     DIRTY,                    /* to count # of dirty segments */
```

```
241     PRE,                      /* to count # of entirely obsolete segments */
```

```
242     NR_DIRTY_TYPE
```

```
243 };
```

```
244
```

245 struct dirty_seglist_info {

```
246     const struct victim_selection *v_ops; /* victim selection operation */
```

```
247     unsigned long *dirty_segmap[NR_DIRTY_TYPE];
```

```
248     struct mutex seglist_lock;          /* lock for segment bitmaps */
```

```
249     int nr_dirty[NR_DIRTY_TYPE];        /* # of dirty segments */
```

```
250     unsigned long *victim_secmap;       /* background GC victims */
```

```
251 };
```

```
252
```

```
253 /* victim selection function for cleaning and SSR */
```

254 struct victim_selection {

```
255     int (*get_victim)(struct f2fs_sb_info *, unsigned int *,
```

```
256                     int, int, char);
```

```
257 };
```

```
258
```

```
259 /* for active log information */
```

260 struct curseg_info {

```
261     struct mutex curseg_mutex;          /* lock for consistency */
```

```
262     struct f2fs_summary_block *sum_blk; /* cached summary block */
```

```
263     struct rw_semaphore journal_rwsem; /* protect journal area */
```

```
264     struct f2fs_journal *journal;       /* cached journal info */
```

```
265     unsigned char alloc_type;           /* current allocation type */
```

```
266     unsigned int segno;                 /* current segment number */
```

```
267     unsigned short next_blkoff;         /* next block offset to write */
```

```
268     unsigned int zone;                  /* current zone number */
```

```
269     unsigned int next_segno;            /* preallocated segment */
```

```
270 };
```

```
271
```

272 struct sit_entry_set {

```
273     struct list_head set_list;          /* link with all sit sets */
```

```
274     unsigned int start_segno;           /* start segno of sits in set */
```

```
275     unsigned int entry_cnt;             /* the # of sit entries in set */
```

```
276 };
```